

Backtesting.py Quick Start User Guide

This tutorial shows some of the features of *backtesting.py*, a Python framework for [backtesting](#) trading strategies.

Backtesting.py is a small and lightweight, blazing fast backtesting framework that uses state-of-the-art Python structures and procedures (Python 3.6+, Pandas, NumPy, Bokeh). It has a very small and simple API that is easy to remember and quickly shape towards meaningful results. The library *doesn't* really support stock picking or trading strategies that rely on arbitrage or multi-asset portfolio rebalancing; instead, it works with an individual tradeable asset at a time and is best suited for optimizing position entrance and exit signal strategies, decisions upon values of technical indicators, and it's also a versatile interactive trade visualization and statistics tool.

Data

You bring your own data. Backtesting ingests *all kinds* of [OHLC](#) data (stocks, forex, futures, crypto, ...) as a [pandas.DataFrame](#) with columns `'Open'`, `'High'`, `'Low'`, `'Close'` and (optionally) `'Volume'`. Such data is widely obtainable, e.g. with packages:

- [pandas-datareader](#),
- [Quandl](#),
- [findatapy](#),
- [yFinance](#),
- [investpy](#), etc.

Besides these columns, **your data frames can have additional columns which are accessible in your strategies in a similar manner**.

DataFrame should ideally be indexed with a *datetime index* (convert it with [pd.to_datetime\(\)](#)); otherwise a simple range index will do.

In [1]:

```
# Example OHLC daily data for Google Inc. from backtesting.test import GOOG GOOG.tail()
```

Loading BokehJS ...

Out[1]:

| | Open | High | Low | Close | Volume |
|------------|--------|--------|--------|--------|---------|
| 2013-02-25 | 802.30 | 808.41 | 790.49 | 790.77 | 2303900 |
| 2013-02-26 | 795.00 | 795.95 | 784.40 | 790.13 | 2202500 |
| 2013-02-27 | 794.80 | 804.75 | 791.11 | 799.78 | 2026100 |
| 2013-02-28 | 801.10 | 806.99 | 801.03 | 801.20 | 2265800 |
| 2013-03-01 | 797.80 | 807.14 | 796.15 | 806.19 | 2175400 |

Strategy

Let's create our first strategy to backtest on these Google data, a simple [moving average \(MA\) cross-over strategy](#).

Backtesting.py doesn't ship its own set of *technical analysis indicators*. Users favoring TA should probably refer to functions from proven indicator libraries, such as [TA-Lib](#) or [Tulipy](#), but for this example, we can define a simple helper moving average function ourselves:

In [2]:

```
import pandas as pd def SMA(values, n): """ Return simple moving average of `values`, at each step taking into account `n` previous values. """ return pd.Series(values).rolling(n).mean()
```

A new strategy needs to extend [Strategy](#) class and override its two abstract methods: [init\(\)](#) and [next\(\)](#).

Method [init\(\)](#) is invoked before the strategy is run. Within it, one ideally precomputes in efficient, vectorized manner whatever indicators and signals the strategy depends on.

Method [next\(\)](#) is then iteratively called by the [Backtest](#) instance, once for each data point (data frame row), simulating the incremental availability of each new full candlestick bar.

Note, *backtesting.py* cannot make decisions / trades *within* candlesticks — any new orders are executed on the next candle's *open* (or the current candle's *close* if [trade_on_close=True](#)). If you find yourself wishing to trade within candlesticks (e.g. daytrading), you instead need to begin with more fine-grained (e.g. hourly) data.

In [3]:

```
from backtesting import Strategy from backtesting.lib import crossover class SmaCross(Strategy): # Define the two MA lags as *class variables* # for later optimization n1 = 10 n2 = 20 def init(self): # Precompute the two moving averages self.sma1 = self.I(SMA, self.data.Close, self.n1) self.sma2 = self.I(SMA, self.data.Close, self.n2) def next(self): # If sma1 crosses above sma2, close any existing # short trades, and buy the asset if crossover(self.sma1, self.sma2): self.position.close() self.buy() # Else, if sma1 crosses below sma2, close any existing # long trades, and sell the asset elif crossover(self.sma2, self.sma1): self.position.close() self.sell()
```

In [init\(\)](#) as well as in [next\(\)](#), the data the strategy is simulated on is available as an instance variable [self.data](#).

In [init\(\)](#), we declare and **compute indicators indirectly by wrapping them in [self.I\(\)](#)**. The wrapper is passed a function (our `SMA` function) along with any arguments to call it with (our *close* values and the MA lag). Indicators wrapped in this way will be automatically plotted, and their legend strings will be intelligently inferred.

In [next\(\)](#), we simply check if the faster moving average just crossed over the slower one. If it did upwards, we close the possible short position and go long; if it did downwards, we close the open long position and go short. Note, we don't adjust order size, so *Backtesting.py* assumes *maximal possible position*. We use [backtesting.lib.crossover\(\)](#) function instead of writing more obscure and confusing conditions, such as:

In [4]:

```
%%script echo def next(self): if (self.sma1[-2] < self.sma2[-2] and self.sma1[-1] > self.sma2[-1]): self.position.close() self.buy() elif (self.sma1[-2] > self.sma2[-2] and # Ugh! self.sma1[-1] < self.sma2[-1]): self.position.close() self.sell()
```

In [init\(\)](#), the whole series of points was available, whereas in [next\(\)](#), **the length of [self.data](#) and all declared indicators is adjusted** on each [next\(\)](#) call so that `array[-1]` (e.g. `self.data.Close[-1]` or `self.sma1[-1]`) always contains the most recent value, `array[-2]` the previous value, etc. (ordinary Python indexing of ascending-sorted 1D arrays).

Note: `self.data` and any indicators wrapped with `self.I` (e.g. `self.sma1`) are NumPy arrays for performance reasons. If you prefer pandas Series or DataFrame objects, use `Strategy.data.<column>.s` or `Strategy.data.df` accessors respectively. You could also construct the series manually, e.g. `pd.Series(self.data.Close, index=self.data.index)`.

We might avoid `self.position.close()` calls if we primed the [Backtest](#) instance with `Backtest(..., exclusive_orders=True)`.

Backtesting

Let's see how our strategy performs on historical Google data. The `Backtest` instance is initialized with OHLC data and a strategy *class* (see API reference for additional options), and we begin with 10,000 units of cash and set broker's commission to realistic 0.2%.

In [5]:

```
from backtesting import Backtest bt = Backtest(GOOG, SmaCross, cash=10_000, commission=.002) stats = bt.run() stats
```

Out[5]:

```
Start 2004-08-19 00:00:00 End 2013-03-01 00:00:00 Duration 3116 days 00:00:00 Exposure Time [%] 94.27
Equity Final [$] 56263.52 Equity Peak [$] 56309.06 Commissions [$] 10563.95 Return [%] 462.64 Buy & Hold
Return [%] 607.37 Return (Ann.) [%] 22.47 Volatility (Ann.) [%] 37.41 CAGR [%] 14.99 Sharpe Ratio 0.60
Sortino Ratio 1.14 Calmar Ratio 0.66 Alpha [%] 450.62 Beta 0.02 Max. Drawdown [%] -33.93 Avg. Drawdown
[%] -6.16 Max. Drawdown Duration 830 days 00:00:00 Avg. Drawdown Duration 50 days 00:00:00 # Trades 93
Win Rate [%] 52.69 Best Trade [%] 56.92 Worst Trade [%] -16.83 Avg. Trade [%] 1.76 Max. Trade Duration
121 days 00:00:00 Avg. Trade Duration 32 days 00:00:00 Profit Factor 1.99 Expectancy [%] 2.29 SQN 1.58
Kelly Criterion 0.21 _strategy SmaCross _equity_curve Equ... _trades Size EntryB... dtype: object
```

`Backtest.run()` method returns a pandas Series of simulation results and statistics associated with our strategy. We see that this simple strategy makes almost 600% return in the period of 9 years, with maximum drawdown 33%, and with longest drawdown period spanning almost two years ...

`Backtest.plot()` method provides the same insights in a more visual form.

In [6]:

```
bt.plot()
```

Out[6]:

```
GridPlot(id = 'p1349', ...)
```

Optimization

We hard-coded the two lag parameters (`n1` and `n2`) into our strategy above. However, the strategy may work better with 15–30 or some other cross-over. **We declared the parameters as optimizable by making them [class variables](#).**

We optimize the two parameters by calling `Backtest.optimize()` method with each parameter a keyword argument pointing to its pool of possible values to test. Parameter `n1` is tested for values in range between 5 and 30 and parameter `n2` for values between 10 and 70, respectively. Some combinations of values of the two parameters are invalid, i.e. `n1` should not be *larger than* or equal to `n2`. We limit admissible parameter combinations with an *ad hoc* constraint function, which takes in the parameters and returns `True` (i.e. admissible) whenever `n1` is less than `n2`. Additionally, we search for such parameter combination that maximizes return over the observed period. We could instead choose to optimize any other key from the returned `stats` series.

In [7]:

```
%%time stats = bt.optimize(n1=range(5, 30, 5), n2=range(10, 70, 5), maximize='Equity Final [$]', constraint=lambda param: param.n1 < param.n2) stats
```

```
CPU times: user 128 ms, sys: 43.3 ms, total: 171 ms Wall time: 1.66 s
```

Out[7]:

```
Start 2004-08-19 00:00:00 End 2013-03-01 00:00:00 Duration 3116 days 00:00:00 Exposure Time [%] 98.14
Equity Final [$] 77829.05 Equity Peak [$] 84982.19 Commissions [$] 30771.04 Return [%] 678.29 Buy & Hold
Return [%] 687.99 Return (Ann.) [%] 27.22 Volatility (Ann.) [%] 43.21 CAGR [%] 18.05 Sharpe Ratio 0.63
Sortino Ratio 1.28 Calmar Ratio 0.61 Alpha [%] 614.80 Beta 0.09 Max. Drawdown [%] -44.55 Avg. Drawdown
[%] -5.81 Max. Drawdown Duration 1558 days 00:00:00 Avg. Drawdown Duration 50 days 00:00:00 # Trades 152
Win Rate [%] 50.00 Best Trade [%] 61.36 Worst Trade [%] -19.98 Avg. Trade [%] 1.32 Max. Trade Duration
83 days 00:00:00 Avg. Trade Duration 21 days 00:00:00 Profit Factor 1.83 Expectancy [%] 1.75 SQN 1.28
Kelly Criterion 0.13 _strategy SmaCross(n1=10,n... _equity_curve Equ... _trades Size Entry... dtype:
object
```

We can look into `stats['_strategy']` to access the Strategy *instance* and its optimal parameter values (10 and 15).

In [8]:

```
stats._strategy
```

Out[8]:

```
<Strategy SmaCross(n1=10,n2=15)>
```

In [9]:

```
bt.plot(plot_volume=False, plot_pl=False)
```

Out[9]:

```
GridPlot(id = 'p1618', ...)
```

Strategy optimization managed to up its initial performance *on in-sample data* by almost 50% and even beat simple [buy & hold](#). In real life optimization, however, do [take steps to avoid overfitting](#).

Trade data

In addition to backtest statistics returned by [`Backtest.run\(\)`](#) shown above, you can look into *individual trade returns* and the changing *equity curve* and *drawdown* by inspecting the last few, internal keys in the result series.

In [10]:

```
stats.tail()
```

Out[10]:

```
SQN 1.28 Kelly Criterion 0.13 _strategy SmaCross(n1=10,n2=15) _equity_curve Equity DrawdownPct
DrawdownDurat... _trades Size EntryBar ExitBar EntryPrice Exit... dtype: object
```

The columns should be self-explanatory.

In [11]:

```
stats['_equity_curve'] # Contains equity/drawdown curves. DrawdownDuration is only defined at ends of DD
periods.
```

Out[11]:

| | | Equity | DrawdownPct | DrawdownDuration |
|------------|--|----------|-------------|------------------|
| 2004-08-19 | | 10000.00 | 0.00 | NaT |
| 2004-08-20 | | 10000.00 | 0.00 | NaT |
| 2004-08-23 | | 10000.00 | 0.00 | NaT |
| 2004-08-24 | | 10000.00 | 0.00 | NaT |
| 2004-08-25 | | 10000.00 | 0.00 | NaT |
| ... | | ... | ... | ... |
| 2013-02-25 | | 76348.73 | 0.10 | NaT |
| 2013-02-26 | | 76287.29 | 0.10 | NaT |
| 2013-02-27 | | 77213.69 | 0.09 | NaT |
| 2013-02-28 | | 77350.01 | 0.09 | NaT |
| 2013-03-01 | | 77829.05 | 0.08 | 1558 days |

2148 rows × 3 columns

In [12]:

```
stats['_trades'] # Contains individual trade data
```

Out[12]:

| | Size | EntryB | ExitB | EntryPrice | PriceSL | TP | P <small>l</small> | Commission | Symbol | EntryTime | ExitTime | Iteration | Strategy | SMA (SMA1) | SMA (SMA2) | SMA (SMA3) | SMA (C,15) | | |
|-----|------|--------|-------|------------|---------|------|--------------------|------------|--------|-----------|------------|------------|----------|------------|------------|------------|------------|--------|--------|
| 0 | 87 | 20 | 60 | 114.42 | 185.23 | None | None | 6108.33 | 32.14 | 0.61 | 2004-08-19 | 2004-08-20 | 56 days | 11-12 | None | 107.40 | 181.07 | 105.73 | 183.32 |
| 1 | -86 | 60 | 69 | 185.23 | 175.80 | None | None | 748.88 | 62.10 | 0.05 | 2004-08-21 | 2004-08-22 | 14 days | 11-26 | None | 181.07 | 173.56 | 183.32 | 173.14 |
| 2 | 95 | 69 | 71 | 175.80 | 180.71 | None | None | 398.71 | 67.74 | 0.02 | 2004-08-21 | 2004-08-22 | 4 days | 11-30 | None | 173.56 | 173.18 | 173.14 | 174.55 |
| 3 | -95 | 71 | 75 | 180.71 | 179.13 | None | None | 81.73 | 68.37 | 0.00 | 2004-08-21 | 2004-08-22 | 6 days | 12-06 | None | 173.18 | 176.58 | 174.55 | 175.51 |
| 4 | 96 | 75 | 82 | 179.13 | 177.99 | None | None | -178.06 | 68.57 | -0.01 | 2004-08-21 | 2004-08-22 | 9 days | 12-15 | None | 176.58 | 175.18 | 175.51 | 176.58 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | | |
| 147 | -90 | 2056 | 2085 | 740.13 | 687.78 | None | None | 4454.48 | 57.02 | 0.07 | 2012-01-01 | 2012-01-10 | 44 days | 11-29 | None | 752.66 | 667.39 | 753.99 | 664.45 |
| 148 | 104 | 2085 | 2111 | 687.78 | 735.54 | None | None | 4670.99 | 96.05 | 0.07 | 2012-01-01 | 2012-01-08 | 40 days | 1-08 | None | 667.39 | 718.50 | 664.45 | 719.00 |
| 149 | -103 | 2111 | 2113 | 735.54 | 742.83 | None | None | -1055.30 | 4.54 | 0.01 | 2013-01-01 | 2013-01-08 | 2 days | 1-10 | None | 718.50 | 724.67 | 719.00 | 721.51 |
| 150 | 101 | 2113 | 2121 | 742.83 | 735.99 | None | None | -989.52 | 98.72 | 0.01 | 2013-01-01 | 2013-01-23 | 13 days | 1-23 | None | 724.67 | 724.37 | 721.57 | 726.41 |
| 151 | -100 | 2121 | 2127 | 735.99 | 750.51 | None | None | -1749.29 | 7.30 | 0.02 | 2013-01-01 | 2013-01-31 | 8 days | 1-31 | None | 724.37 | 738.20 | 726.47 | 735.12 |

152 rows × 18 columns

Learn more by exploring further [examples](#) or find more framework options in the [full API reference](#).